

Live Lab: Apply Structured Prompt Templates

Scenario

In this lab you will learn how structured prompt templates can standardize AI outputs, reduce ambiguity, and enforce professional formats in security operations. By working with firewall logs, you will see how predefined templates guide analysis and reporting, making it easier to compare results across analysts. You will also practice completing partially filled templates, correcting flawed examples, and adapting templates for your own needs.

Your Mission

- Apply structured prompt templates to real security log data.
- Compare freeform prompts against structured templates.
- Fill in and customize template fields to enforce consistency.
- Use helper scripts to simplify structured templates.

5.3.1 Security Alert Summary Template

Pre-requisite step: Download the files in Labfiles5.3 from the GitHub repository. Move it to /home/kali/Labfiles5.3

Step 1: From the desktop, open **Firefox**. Open and sign into **Open WebUI**.

Step 2: From the New Chat page, select the model dropdown menu and select the **mistral:7b** model.

Step 3: In the chat, enter the following prompt and note the results:

```
Prompt 1:
You are a SOC analyst. Read the INCIDENT CARD and write a concise
paragraph describing what happened in plain language. Do not add facts not
present in the card.
```

```
INCIDENT CARD (PHISHING - FINANCE)
- Date (CT): 2025-09-14
- First observed: 09:12:33
- Last observed: 10:05:11
- Targets: Finance staff (~25 inboxes)
- User reports: 3 (two clicked, one entered credentials is unconfirmed)
- Subject (examples): "Payment Advice - URGENT", "Invoice Past Due"
- From: accounts@billingsafe.co (Return-Path: bounce@billingsafe.co)
- Domains: billingsafe.co (registered 2025-09-10)
- Auth results: DMARC=fail, SPF=softfail, DKIM=none
- Links: https://invoice-billingsafe[.]co/login (Microsoft 365 lookalike)
- Attachments: PaymentAdvice.htm (HTML)
- Controls: Initial quarantine applied; gateway rule updated; some emails
reached users
- Known impact: None confirmed; credential theft possible but unverified
```

Insights: Asking an open-ended question gives you a baseline. You see what the model produces without guidance — its natural level of detail, structure, and terminology. This makes it easier to appreciate the improvements once you apply a template, because you'll have something unstructured to compare against.

Step 4: In the same chat, enter the following prompt and note the response:

Prompt 2:

Complete the SECURITY ALERT SUMMARY TEMPLATE strictly from the INCIDENT CARD. If a field is unknown, write UNKNOWN.

SECURITY ALERT SUMMARY TEMPLATE

THREAT CLASSIFICATION:

- Severity Level: [CRITICAL/HIGH/MEDIUM/LOW]
- Threat Type: Phishing (Credential Harvesting)
- Confidence Level: [HIGH/MEDIUM/LOW]

TECHNICAL SUMMARY:

- Source: [sender/return-path]
- Target: [teams/departments]
- Attack Vector: Email, Link to fake login
- Affected Services: Email, Authentication

INDICATORS OF COMPROMISE (IOCs):

- Sender Domain(s): [list]
- URL(s): [list]
- Auth Results: [DMARC/SPF/DKIM]
- Attachment(s): [name/type]

TIMELINE:

- First Observed: [time]
- Last Observed: [time]
- Duration: [compute or UNKNOWN]

IMPACT ASSESSMENT:

- Systems Affected: [range or UNKNOWN]
- Data at Risk: Credentials (POSSIBLE)
- Service Disruption: Minimal

RECOMMENDED ACTIONS:

- 1.
- 2.
- 3.

NEXT STEPS:

- 1.
- 2.
- 3.

Insights: This template is designed for SOC analysts and incident responders. It standardizes how security events are documented so that every alert includes key elements: classification, technical details, indicators of compromise, timeline, impact, and recommended actions. By

enforcing this structure, the template ensures that alerts are complete, consistent, and actionable, regardless of which analyst prepares them.

5.3.2 Incident Triage Template

When multiple alerts come in, analysts need a fast way to sort the noise from the real threats. The Incident Triage template helps capture the essential details right away so incidents can be classified, prioritized, and escalated if necessary.

Step 1: In the same mistral chat, enter the following prompt and note the response:

Prompt 3:

Complete the INCIDENT TRIAGE ASSESSMENT strictly from the INCIDENT CARD. Use the rubric: P2-HIGH if multiple users affected and credential theft is plausible but unconfirmed.

INCIDENT TRIAGE ASSESSMENT

INCIDENT METADATA:

- Incident ID: INC-2025-0914-FIN-01
- Date/Time Reported: 2025-09-14T10:15:00 CT
- Reporter: SOC queue (user reports escalated)
- Initial Source: Email Security Alerts + User Reports

INCIDENT CLASSIFICATION:

- Category:
- Subcategory:
- Severity: [P1/P2/P3/P4]
- Urgency: [IMMEDIATE/4HR/24HR/48HR]

SCOPE ASSESSMENT:

- Affected Systems:
- Affected Users:
- Affected Services:
- Geographic Scope:

TECHNICAL INDICATORS:

- Attack Vectors:
- Suspicious Domains:
- URLs:
- Attachments:
- Auth Results:

BUSINESS IMPACT:

- Service Availability: Operational
- Data Confidentiality: POTENTIALLY_BREACHED (credentials)
- System Integrity: Intact
- Regulatory Implications: UNKNOWN

Insights: This template is used by SOC analysts and shift leads during the first pass of an investigation. Its purpose is to quickly classify an incident, estimate scope, assign severity and urgency, and decide whether it needs to be escalated. By capturing this information in a consistent format, teams can prioritize incidents effectively and ensure the most critical threats are handled first.

5.3.3 Executive Summary Template

Technical details are vital for analysts, but executives need information framed in terms of business impact and decisions. The Executive Summary template condenses complex incidents into plain language, highlighting only the most critical findings, risks, and recommended actions.

Step 1: In the same mistral chat, enter the following prompt and note the response:

```
Prompt 4:
Produce a CYBERSECURITY EXECUTIVE SUMMARY for leadership, based strictly
on the INCIDENT CARD.

CYBERSECURITY EXECUTIVE SUMMARY

EXECUTIVE OVERVIEW:
[2 sentences, plain language, no acronyms.]

KEY METRICS:
- Incidents Detected:
- Incidents Resolved:
- Mean Time to Detection:
- Mean Time to Response:
- Systems Affected:

BUSINESS IMPACT:
- Immediate:
- Potential:
- Financial:

ACTIONS TAKEN:
- Quarantined messages and tightened gateway rules
- Blocked sender domains/URLs
- User notifications initiated

RECOMMENDATIONS:
1.
2.
3.
```

Insights: This template is written for executives and senior decision-makers. It strips away technical jargon and focuses on business impact, risk, and next steps. By presenting incidents in a concise, plain-language format, it ensures leadership can make informed decisions quickly without needing to parse detailed logs or analyst notes.

5.3.4 Prompt Templates with Internal Data

Now that you've practiced templates on general scenarios, it's time to apply one to real security data. In this step, you'll upload an alert file and use a structured prompt template to generate a standardized alert summary from actual network activity.

Step 1: In Open WebUI, start a new chat with **mistral**.

Step 2: Beneath the chat box, select **+ More** and then select **Upload Files**.

Step 3: Select **Alert File1.json** from **/home/kali/Labfiles5.3**, then select **Open**.

Step 4: In the chat, enter the following prompt and note the response:

```
Prompt 5:
Analyze the json file, output:
1) Summary (2 sentences) using info, analysis, threat_level_id, tags.
2) Key IOCs (max 5 bullets).
3) Recommendations (3 short bullets).

If something is missing, write UNKNOWN. Do not add facts.
```

: This will run the prompt against the uploaded file.

Insights: When you upload a large file straight into the model, it has to read the whole dataset in one go. This makes responses slow, can cause the model to skip or oversimplify details, and often leads to inconsistent outputs. In practice, security teams use Retrieval-Augmented Generation (RAG): files are broken into indexed chunks, and the model only pulls in the most relevant pieces. RAG makes analysis faster, more accurate, and scalable, which you'll explore in a later lab.

5.3.5 Templates with Helper Scripts

Now that you've seen how direct uploads work, let's try a different approach. We'll use a Python helper script to preprocess the logs and feed the model a clean, structured prompt. This makes the workflow faster and the outputs more reliable.

Step 1: From the left menu, open the **Terminal**.

Step 2: In the terminal, run the following command:

```
cd /home/kali/Labfiles5.3
python3 security_template1.py
```

Insights: The Alert File1.json raw data is processed by the Python helper script, which builds a prompt from a template and sends it to the model for analysis. Rather than asking the model to sift through all of the raw JSON, the script extracts the relevant fields, structures them, and passes only what's needed. This approach unlocks the power of Python scripting — you can preprocess data, fill in template variables automatically, enforce formats, and even generate different prompts depending on thresholds or filters.

Step 3: Wait a minute or two for the model response to appear.

Insights: The helper script talks directly to the model backend through its API. You run it in the terminal, but behind the scenes it sends a structured prompt and collects the model's response.

Step 4: In the terminal, run the following command:

```
python3 security_template2.py
```

Insights: This script works a bit differently from the first one. It runs in an interactive chat mode where responses are streamed in real time, so you can watch the model generate its answer. It also supports multi-turn conversations, meaning you can reply to the model's initial output and continue the dialogue about the analyzed file.

Step 5: Once the model output completes, enter the following prompt in response:

Prompt 6:
What can you tell me about Clipbanker?

Insights: Enabling multi-turn chat allows you to ask follow-up questions in the terminal session, providing further clarification and detail for investigation or reporting.

5.3.6 Templates with Open WebUI Tools

You've seen how the Python helper script can analyze the JSON alert file from the terminal, with support for multi-turn conversations. Now, let's look at how that same logic has been integrated directly into Open WebUI as a custom Python tool.

Step 1: Close the terminal and return to **Open WebUI**.

Step 2: In Open WebUI, start a new chat with **gemini or any free model from OpenRouter**.

Step 3: Beneath the chat box, select **+ More** and then select **Upload Files**.

Step 4: Select **Alert File1.json** from **/home/kali/Labfiles5.3**, then select **Open**.

Step 5: In the chat, enter the following prompt and note the response:

```
Prompt 7:  
Analyze the uploaded alert file and provide a structured summary.
```

Insights: How this works:

- The tool code wraps the helper script logic into a Python function that processes the uploaded file.
- The tool loads the JSON, extracts metadata, and builds a structured prompt for the model.
- The model never directly “opens” the file. Instead, it receives a carefully crafted text summary of the contents.
- This is a bridge approach — not as scalable as RAG, but still effective for one-off analysis tasks.

Step 6: Verify prompts by selecting the verify button in the lab interface.

⚠ **Security Note:** The tools used in this lab inherently enable remote code execution (RCE). That capability is powerful and valuable, but it also underscores the importance of using only trusted tools and sources. Even when execution is deliberate and beneficial, you must ensure the code being run comes from secure, verified origins and that the level of access granted is appropriate for your environment.

By extending the helper script into an Open WebUI tool, you've seen how structured prompt templates can be operationalized beyond the terminal. The same template logic is now

accessible directly in chat, where it benefits from consistent formatting, repeatable analysis, and seamless integration with your AI workflows.